

Introduction to Bullshit: From Philosophy to Practice

A Critical Thinking Journey

Brendan Shea, PhD

Intro to Logic

What Is Bullshit? An Introduction to a Philosophical Problem

- **Bullshit** is a distinct form of communication that differs from both truth-telling and lying in important ways.
- Unlike a liar who knows the truth and deliberately says something false, a bullshitter shows complete indifference to whether their statements are true or false.
- The philosophical study of bullshit helps us understand how language can be used to manipulate and obscure rather than communicate.
- Recognizing bullshit is a crucial skill in our information-rich world where not all speech aims at truth.

Key Insight

Bullshit is more dangerous than lies because it erodes our very concern for truth.

Why Study Bullshit in a Logic Class?

- Logic teaches us to evaluate arguments based on their structure and the truth of their premises, but bullshit operates outside this framework.
- Understanding bullshit strengthens our **critical thinking** skills by teaching us to recognize when someone isn't even trying to make valid arguments.
- In a world full of social media, AI-generated content, and political rhetoric, the ability to detect bullshit has become essential for informed citizenship.
- Studying bullshit connects abstract logical principles to real-world communication challenges you face every day.

Connection to Logic

Traditional logic assumes speakers aim for truth. Bullshit analysis examines what happens when this assumption breaks down.

Overview: From Theory to Real-World Applications

- We'll begin with philosopher Harry Frankfurt's groundbreaking theory that defines bullshit as speech unconcerned with truth or falsity.
- Next, we'll examine how social media platforms create perfect conditions for bullshit to thrive and spread rapidly.
- We'll then explore how artificial intelligence can generate plausible-sounding but potentially meaningless or false content at unprecedented scale.
- Finally, we'll analyze how political bullshit, especially from populist movements, can threaten the shared reality necessary for democratic deliberation.

What we'll cover

- 1 Frankfurt's philosophical framework
- 2 Social media case studies
- 3 AI and synthetic content
- 4 Political applications and democracy

Harry Frankfurt: The Philosopher Who Defined Bullshit


- Harry Frankfurt is a distinguished philosopher at Princeton University who spent most of his career studying free will, personal identity, and moral responsibility.
- In 1986, he published a short essay called "On Bullshit" that became his most famous work, later republished as a bestselling book in 2005.
- Frankfurt argued that bullshit deserved serious philosophical analysis because it represents a unique and dangerous form of dishonesty.
- His work transformed bullshit from a vulgar term into a **technical concept** in philosophy, communication studies, and critical thinking.

Frankfurt's Contribution

"One of the most salient features of our culture is that there is so much bullshit. Everyone knows this." - Opening line of "On Bullshit"

"On Bullshit": The Essay That Started It All

- Frankfurt's essay begins by noting that despite bullshit's prevalence in our culture, we lack a clear understanding of what it actually is.
- He argues that **bullshit** is not merely casual talk or hot air, but a specific form of deceptive communication with its own characteristics.
- The essay uses philosophical analysis to distinguish bullshit from related concepts like lying, bluffing, and mere carelessness with facts.
- Frankfurt's central thesis is that the **essence of bullshit** is a complete lack of concern for how things really are—an indifference to truth itself.



Truth, Lies, and the Space Between

- When someone tells the **truth**, they attempt to describe reality accurately based on their genuine beliefs about how things are.
- When someone tells a **lie**, they know what is true but deliberately assert something false to deceive their audience.
- **Bullshit** occupies a third space: the speaker neither knows nor cares whether their statements correspond to reality.
- This indifference to truth makes bullshit fundamentally different from both honest mistakes and calculated deceptions.

Type	Knows Truth?	Cares About Truth?
Truth-teller	Yes/Tries to	Yes
Liar	Yes	Yes (to violate)
Bullshitter	No/Maybe	No

The Bullshitter's Relationship to Truth

- The bullshitter's **indifference to truth** is what makes bullshit uniquely corrosive to honest discourse and rational deliberation.
- While a liar must pay attention to the truth in order to conceal it effectively, a bullshitter ignores the truth entirely.
- Bullshitters may sometimes say true things and sometimes say false things, but truth or falsity is never their goal or concern.
- This disconnection from truth-seeking makes bullshit more dangerous than lying because it attacks the very enterprise of trying to establish what is real.

Frankfurt's Warning

The bullshitter "does not care whether the things he says describe reality correctly. He just picks them out, or makes them up, to suit his purpose."

Indifference vs. Deception: The Key Distinction

- **Deception** requires the deceiver to have a specific false belief they want their audience to adopt, which means they must track what is actually true.
- **Indifference** means the speaker doesn't care whether their audience's resulting beliefs are true or false—only that they achieve some other goal.
- A liar says "It's sunny outside" knowing it's raining because they want you to leave without an umbrella, while a bullshitter might say the same thing without even checking the weather.
- This distinction explains why fact-checking often fails against bullshit: the bullshitter isn't trying to avoid being caught in falsehoods.

Example

Liar: "I have a PhD from Harvard" (knows they don't)

Bullshitter: "A Harvard professor said..." (doesn't care whether this is true)

Why Bullshit Might Be Worse Than Lying

- Frankfurt argues that **bullshit poses a greater threat** to truth than lying because it undermines the very concept that truth matters.
- Each lie is a localized assault on truth about a specific fact, while bullshit erodes the general habit of caring whether our statements correspond to reality.
- When bullshit becomes normalized, people stop expecting speakers to even attempt accuracy, which destroys the foundation of rational discourse.
- A culture saturated with bullshit loses the shared commitment to truth that makes productive disagreement and learning possible.

The Greater Danger

- **Lies:** Attack specific truths
- **Bullshit:** Attacks the value of truth itself
- **Result:** Epistemic nihilism—nothing matters

The Bullshitter's Goals: Impression Over Accuracy

- The bullshitter's primary goal is to create a certain **impression of themselves** rather than to communicate information about the world.
- Common bullshitter goals include appearing knowledgeable, seeming important, avoiding blame, or filling awkward silence with confident-sounding words.
- A job applicant who claims expertise in "quantum blockchain synergy" isn't trying to deceive about a real skill—they're trying to sound technically sophisticated.
- This focus on **impression management** explains why bullshitters often use vague buzzwords, complex jargon, and unfalsifiable claims.

Classic Bullshit Phrases

"Studies show..."	(which studies?)
"Everyone knows..."	(who exactly?)
"It's common sense..."	(is it though?)
"Trust me..."	(based on what?)

Frankfurt's Examples: From Art to Politics

- Frankfurt illustrates bullshit with the example of a Fourth of July orator who makes grand pronouncements about patriotism without caring if they're true.
- He discusses how **bull sessions**—casual conversations where people try out ideas without commitment—can slide into bullshit when speakers pretend their musings are serious claims.
- In art and advertising, bullshitters make sweeping statements about products or artworks ("This painting captures the essence of humanity") without concern for accuracy.

Example

Fourth of July Orator: "Our founding fathers would be proud of how we've fulfilled their vision!"

- Never studied what the founders actually wanted
- Doesn't care if statement is historically accurate
- Only cares that audience feels patriotic

The Perfect Storm: Why Social Media Breeds Bullshit


- Social media platforms create ideal conditions for bullshit by rewarding **engagement** (likes, shares, comments) rather than accuracy or truth.
- The pressure to post constantly means users often share content without checking facts, embodying Frankfurt's "indifference to truth."
- The **viral nature** of social media means bullshit can spread to millions before any fact-checking occurs, if it ever does.
- Platform algorithms amplify emotionally provocative content regardless of accuracy, giving bullshit a systematic advantage over careful, truthful communication.

The Engagement Trap

A false but shocking post gets 10,000 shares. The correction gets 100. Which one shaped more people's beliefs?

Algorithmic Amplification of Engaging Content

- Social media **algorithms** are designed to maximize user engagement by showing content that triggers strong emotional responses.
- These algorithms don't evaluate truth or falsity—they only measure clicks, time spent, and interactions, perfectly embodying Frankfurt's "indifference to truth."
- Bullshit often outperforms truth in engagement metrics because it can be crafted purely to push emotional buttons without the constraints of accuracy.
- The result is a **feedback loop**: successful bullshit gets amplified, teaching content creators that truth is less important than virality.



The Attention Economy: Clicks Over Truth

- In the **attention economy**, human attention is the scarce resource that platforms compete to capture and monetize through advertising.
- Content creators learn that provocative bullshit ("Doctors HATE this one trick!") generates more clicks than nuanced truth.
- The economic incentives actively discourage fact-checking or careful analysis because these slow down content production and reduce sensationalism.
- This system rewards those who are best at crafting engaging bullshit, not those who are most committed to truth or accuracy.

Example

- **Truth:** "New study shows modest correlation between coffee consumption and alertness in specific conditions"
- **Bullshit:** "Scientists SHOCKED: Coffee is basically a SUPERPOWER!"
- Guess which headline gets more clicks?

Echo Chambers and Confirmation Bias

- **Echo chambers** form when algorithms show users content similar to what they've already engaged with, creating closed loops of reinforcing beliefs.
- Within these chambers, **confirmation bias** makes people more likely to share content that supports their existing views without checking its accuracy.
- Bullshit that aligns with group beliefs spreads unchallenged because the social cost of fact-checking your "team's" claims is higher than going along.
- This environment is perfect for bullshitters who can craft messages that feel truthy to specific audiences without any concern for actual truth.

The Truthiness Problem

"Truthiness" (coined by Stephen Colbert): When something feels true based on gut instinct rather than facts. Echo chambers transform bullshit into truthiness.

Viral Misinformation: Case Studies

- The "Momo Challenge" hoax spread worldwide through social media, causing panic about a nonexistent threat to children, shared by parents who never verified it.
- False health claims about miracle cures spread faster than medical professionals can debunk them because hope and fear are more shareable than careful science.
- These cases show how **emotional resonance** combined with **indifference to verification** creates perfect conditions for bullshit epidemics.

Example

Anatomy of Viral Bullshit: The "Dangerous TikTok Trend" Template

- 1 Claim teens are doing something shocking
- 2 Add "Parents need to know!"
- 3 Include unverified anecdotes
- 4 Watch it spread without anyone checking if the trend actually exists

Influencer Culture and Authentic Deception

- **Influencer culture** rewards those who can project an appealing image regardless of whether that image corresponds to reality.
- The performance of "authenticity" becomes more important than actual authenticity, creating what we might call **authentic bullshit**.
- Influencers often promote products they've never used or lifestyles they don't actually live, embodying Frankfurt's indifference to truth.
- Followers absorb this bullshit because the emotional connection to the influencer overrides critical evaluation of their claims.

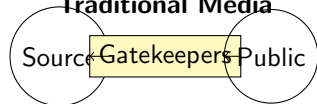
Influencer Bullshit Starter Pack

"I wake up at 4am every day"	(posted at 2pm)
"This changed my life!"	(sponsored post)
"Just being real with you guys"	(reading from script)
"I've used this for years"	(first time touching product)

The Collapse of Gatekeepers

- Traditional media had **gatekeepers**—editors, fact-checkers, and journalists—who filtered out at least some bullshit before publication.
- Social media eliminated these gatekeepers, allowing anyone to broadcast claims to potentially millions without any verification process.
- While democratizing information sounds positive, it also means there's no systematic filter for bullshit before it reaches massive audiences.
- The responsibility for detecting bullshit has shifted entirely to individual users who often lack the time, expertise, or motivation to verify claims.

Traditional Media



Social Media



Digital Literacy: Tools for Detection

- **Digital literacy** means developing skills to navigate online information critically, including recognizing potential bullshit before sharing.
- Key detection strategies include checking sources, looking for emotional manipulation, and asking "Who benefits if I believe this?"
- The **SIFT method** (Stop, Investigate the source, Find better coverage, Trace claims) provides a practical framework for evaluating online content.
- Remember that bullshitters count on you being too busy or lazy to verify—taking just 30 seconds to check can stop bullshit from spreading.

Quick Bullshit Detection Checklist

- 1 Does this seem designed to make me angry or afraid?
- 2 Is the source trying to sell me something?
- 3 Would I be embarrassed if this turned out to be false?
- 4 Can I find this reported by legitimate news sources?

Large Language Models: Bullshit Machines?

- **Large Language Models (LLMs)** like ChatGPT are AI systems trained to produce plausible-sounding text by predicting what words typically follow others.
- These models have no conception of truth or falsity—they simply generate text that statistically resembles their training data.
- This makes LLMs perfect exemplars of Frankfurt's bullshitter: they are fundamentally **indifferent to truth** because they have no mechanism for caring about it.
- When an AI confidently states facts, it's not lying or telling the truth—it's producing text patterns without any relationship to reality.

The Fundamental Problem

LLMs optimize for plausibility, not truth. They literally cannot tell the difference between accurate information and convincing-sounding nonsense.

The Hallucination Problem in AI

- AI researchers use the term **hallucination** when models generate false information that sounds completely reasonable and is stated with total confidence.
- These aren't mistakes in the traditional sense—the AI isn't misremembering or miscalculating; it's generating plausible-sounding bullshit.
- The hallucination problem is inherent to how these models work: they're designed to produce text that fits patterns, not text that corresponds to reality.

Example

Real AI Hallucination Examples:

- Lawyer used ChatGPT for legal research; it invented entire court cases
- Students asked for book summaries; AI created convincing plots for books that never existed
- Researchers requested citations; AI generated real-sounding but completely fake papers

ChatGPT and Confident Incorrectness

- ChatGPT and similar models present all information with the same confident, authoritative tone regardless of accuracy.
- This **uniform confidence** makes it impossible to distinguish between the AI's accurate statements and its complete fabrications.
- The model cannot say "I don't know" in any meaningful way—it will always produce something that sounds like an answer.
- Users often trust AI responses because they're well-formatted and sound professional, mistaking stylistic competence for factual accuracy.

The Confidence Game

Human Expert: "I think... probably... let me check..."

ChatGPT (correct): "The answer is definitively X."

ChatGPT (hallucinating): "The answer is definitively Y."

Same confidence level, opposite truth values

AI-Generated Content: Plausible but False

- AI excels at generating content that **sounds right** because it mimics patterns from millions of real examples in its training data.
- This plausibility makes AI-generated bullshit particularly dangerous—it passes our initial "sniff test" for reasonable-sounding information.
- AI can produce entire essays, news articles, or scientific abstracts that are structurally perfect but factually meaningless or false.
- The sophistication of the bullshit makes it harder to detect than human-generated nonsense, which often contains tell-tale signs of carelessness.

Example

AI-Generated Abstract: "We present a novel framework for quantum-enhanced machine learning using topological insulators. Our results show a 47% improvement in classification accuracy through non-Abelian anyonic braiding..."

- Sounds impressively technical
- Uses real scientific terms
- Completely meaningless combination

Deepfakes and Synthetic Media

- **Deepfakes** use AI to create convincing fake videos or audio of real people saying things they never said.
- This technology represents bullshit at a new level: not just false claims, but fabricated evidence that seems to support those claims.
- The existence of deepfakes creates a **liar's dividend**—now anyone can dismiss real evidence by claiming it might be AI-generated.
- We're entering an era where seeing is no longer believing, making Frankfurt's concern about indifference to truth more urgent than ever.

The Epistemic Apocalypse

When any video could be fake and any fake could look real, how do we establish shared facts about what actually happened?

The Flood of AI-Generated Text Online

- AI tools can generate thousands of articles, reviews, or social media posts per minute, flooding the internet with synthetic content.
- This **bullshit tsunami** makes it increasingly difficult to find genuine human communication or verified information online.
- SEO-optimized AI content often ranks higher in search results than carefully researched human-written articles, rewarding bullshit over truth.
- Soon, most text online may be AI-generated, creating a "wilderness of mirrors" where detecting authentic human communication becomes nearly impossible.

Can AI Detect Its Own Bullshit?

- Researchers are developing AI systems to detect AI-generated content, but this creates an **arms race** between generation and detection.
- Current detection tools have high error rates and can be fooled by sophisticated AI or flag genuine human writing as artificial.
- The fundamental problem remains: if AI doesn't understand truth, how can it identify bullshit based on indifference to truth?
- We may need to accept that technical solutions alone cannot solve what is essentially a human problem about our relationship with truth.

The Detection Paradox

- AI creates convincing bullshit
- We build AI to detect AI bullshit
- That AI also doesn't understand truth
- Who detects the detector's bullshit?

Teaching Critical Thinking in the AI Age

- Traditional critical thinking skills become even more essential when anyone can generate professional-looking bullshit with a few clicks.
- Students need to learn that **polish doesn't equal truth**—AI can make false claims look more credible than messy human truths.
- We must teach the habit of asking "What is the source's relationship to truth?" not just "Does this sound plausible?"
- The goal isn't to reject all AI-assisted content but to maintain healthy skepticism and verify important claims regardless of how authoritative they appear.

New Critical Thinking Questions for the AI Age

- 1 Could this be AI-generated? How would I know?
- 2 What specific facts can I verify independently?
- 3 Why might someone use AI to create this content?
- 4 Am I being fooled by style over substance?

Political Bullshit: A Historical Perspective

- Political bullshit isn't new—Plato worried about sophists who taught rhetoric without regard for truth 2,400 years ago.
- What's changed is the **scale and speed** at which political bullshit can spread through modern media ecosystems.
- Historical examples show how political movements have used bullshit to create alternate realities that justify their power.
- Understanding this history helps us recognize that defending truth against bullshit is an ongoing challenge for every generation.

Example

Timeless Political Bullshit Tactics:

- Ancient Rome: "Bread and circuses" to distract from real issues
- 1930s: "Big Lie" technique—repeat falsehoods until believed
- Cold War: "Whataboutism" to deflect criticism
- Today: "Firehose of falsehood"—overwhelm with volume

Populism and the Appeal of Simple Answers

- **Populist movements** often rely on bullshit by offering simple, emotionally satisfying answers to complex problems.
- The populist bullshitter doesn't care if "the elite" are actually responsible for every problem—they care that this narrative resonates with frustrated voters.
- Complex truths about economic systems, global interdependence, or policy tradeoffs can't compete with simple bullshit that identifies clear villains.
- This creates a **race to the bottom** where politicians who acknowledge complexity lose to those who traffic in compelling oversimplifications.

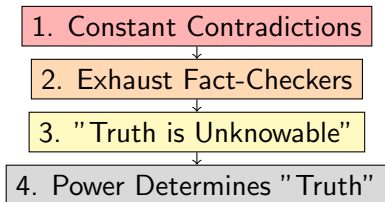
The Simplicity Trap

Real problem: "Declining manufacturing jobs due to automation, globalization, changing markets, education gaps..." Populist bullshit: "They took your jobs!" Guess which message wins elections?

The Authoritarian's Playbook: Flooding the Zone

- Modern authoritarians use a strategy called "**flooding the zone with shit**"—producing so much bullshit that citizens give up on finding truth.
- The goal isn't to convince people of specific lies but to exhaust them into believing that all claims are equally unreliable.
- This strategy weaponizes Frankfurt's insight: destroying the very concept of truth is more powerful than promoting specific falsehoods.
- When people conclude "everyone lies," they stop holding leaders accountable for dishonesty and retreat into cynicism.

Authoritarian Bullshit Strategy



Case Study: Alternative Facts and Post-Truth

- The phrase "**alternative facts**" entered public discourse in 2017, perfectly embodying Frankfurt's concept of indifference to truth.
- This wasn't claiming that specific statements were true or false, but that the very concept of objective facts could be dismissed.
- The term "**post-truth**" describes our current era where emotional appeals and tribal loyalty matter more than factual accuracy.
- These concepts represent the mainstreaming of bullshit—the open acknowledgment that truth is optional in political discourse.

Example

Evolution of Political Communication:

- Traditional lie: "The crowd was 1 million" (knows it was 250,000)
- Bullshit: "It was the biggest crowd ever" (doesn't care about actual number)
- Post-truth: "Crowd size is whatever we say it is" (rejects measurement itself)

How Democracies Become Vulnerable

- **Liberal democracy** depends on citizens sharing enough common facts to have productive disagreements about values and policies.
- When bullshit destroys this shared factual foundation, democratic deliberation becomes impossible—we can't debate solutions if we can't agree on problems.
- Democratic norms like good-faith argument and compromise assume that participants care about truth, an assumption bullshitters exploit.
- Ironically, democracy's protections for free speech make it vulnerable to those who use that freedom to undermine the very possibility of rational discourse.

The Democratic Dilemma

Democracy requires:

- Citizens who can evaluate claims rationally
- Leaders who engage in good faith
- Shared commitment to truth-seeking

Bullshit destroys all three.

The Erosion of Shared Reality

- Political bullshit creates separate **information universes** where different groups literally cannot agree on basic facts about what happened.
- This goes beyond normal political disagreement—it's not arguing about what policies best serve our values, but about whether events occurred at all.
- When each side accuses the other of living in a fantasy world, the concept of objective reality itself becomes politicized.
- Without shared reality as a foundation, political violence becomes more likely as groups see opponents not as wrong but as delusional or evil.

Signs of Eroding Shared Reality

Healthy Democracy

"I disagree with your interpretation"
"Your sources are biased"
"Let's examine the evidence"
"We have different values"


Bullshit-Damaged Democracy

"That never happened"
"All sources are fake"
"Evidence is a conspiracy"
"You live in an alternate reality"

Conspiracy Theories as Political Tools

- **Conspiracy theories** represent weaponized bullshit—unfalsifiable narratives that explain everything while proving nothing.
- Political actors promote conspiracies not because they believe them but because they're useful for mobilizing supporters and dismissing criticism.
- The beauty of conspiracy bullshit is that evidence against the theory becomes evidence of how deep the conspiracy goes.
- This creates loyal followers who see the leader as the only source of truth in a world of lies, perfect for authoritarian movements.

The Conspiracy Bullshit Cycle



Media Literacy and Democratic Resilience

- **Media literacy** education is essential for democratic survival in an age of sophisticated political bullshit and information warfare.
- Citizens need skills to identify reliable sources, recognize manipulation tactics, and understand how their own biases make them vulnerable.
- Building **democratic resilience** means creating institutions and norms that can withstand floods of bullshit without abandoning free speech.
- This includes supporting quality journalism, teaching critical thinking in schools, and developing social norms that value truth-seeking over tribal loyalty.

Building Resilience: Individual and Collective

Individual: Develop personal habits of verification and skepticism

Community: Create spaces for good-faith dialogue across differences

Institutional: Support fact-checking and investigative journalism

Cultural: Celebrate truth-seeking over winning arguments

The Cost of Bullshit to Society

- The **social costs** of widespread bullshit go far beyond individual false beliefs to corrode the foundations of civilized society.
- When bullshit dominates, expertise becomes worthless, scientific progress stalls, and problems go unsolved because we can't agree they exist.
- Trust—the invisible infrastructure of modern society—erodes when we assume everyone is bullshitting, making cooperation and commerce difficult.
- The ultimate cost is **civilizational decline**: societies that can't distinguish truth from bullshit can't adapt, learn, or solve collective challenges.

What We Lose to Bullshit

- Public health (vaccine hesitancy, miracle cures)
- Economic stability (market manipulation, fraud)
- Social cohesion (polarization, violence)
- Human progress (science denial, anti-expertise)

Building Your Bullshit Detection Toolkit

- Your personal **bullshit detection toolkit** combines Frankfurt's philosophical insights with practical skills for navigating modern information environments.
- Always ask: "Does this person care whether what they're saying is true?" If not, you've identified a bullshitter.
- Develop habits of verification, but also recognize when you're being overwhelmed with claims designed to exhaust your fact-checking capacity.
- Remember that fighting bullshit isn't just about personal protection—it's about maintaining the social conditions that make truth-seeking possible.

Example


Your Toolkit Checklist:

- 1 **Philosophical:** Understand Frankfurt's framework
- 2 **Emotional:** Notice when content triggers strong feelings
- 3 **Source evaluation:** Check expertise and motivations
- 4 **Cross-reference:** Verify claims through multiple sources

The Responsibility of Citizens in a Democracy

- In a democracy, each citizen bears **personal responsibility** for maintaining the information ecosystem we all depend on.
- This means not just avoiding spreading bullshit ourselves, but actively supporting truth-seeking institutions and calling out bullshit when we see it.
- We must resist the temptation to use bullshit for "our side" even when opponents do—the damage to shared reality affects everyone.
- Young people especially have the power to shape new norms around truth and bullshit in digital spaces you'll inhabit for decades.

The Citizen's Choice



Final Thoughts: Truth Still Matters

- Despite the flood of bullshit from social media, AI, and politics, **truth still matters** because reality doesn't care about our narratives.
- Climate change, pandemics, and economic forces operate regardless of what we choose to believe about them—bullshit offers no protection from consequences.
- The fact that Frankfurt's essay became a bestseller shows that people hunger for clarity about truth and deception in confusing times.
- Your generation will determine whether we rebuild a culture that values truth or surrender to a post-truth world—the choice is yours to make.

Remember Frankfurt's Core Insight

The bullshitter is a greater enemy of truth than the liar, because the bullshitter attacks the very idea that truth matters.

Don't let them win.